Proceedings Volume of the 7th IET International Smart Cities Symposium,
7th SCS23, 3-5 December 2023, University of Bahrain, Bahrain

Data Poisoning: Issues, Challenges, and Needs

Mohammad Aljanabi[1], Alaa Hamza Omran[2], Maad M. Mijwil[3], Mostafa Abotaleb[4], El-Sayed M. El-kenawy[5], Sahar Yousif Mohammed[6], and Abdelhameed Ibrahim[7]

[1]Department of computer, College of Education, Aliarqai university, Baghdad, Iraq
[2]University of information technology and communications, Baghdad, Iraq
[3]Computer Techniques Engineering Department Baghdad College of Economic Sciences University, Baghdad, Iraq,
[4]Department of System Programming, South Ural State University, Chelyabinsk, Russia, abotalebmostafa@bk.ru
[5]Department of Communications and Electronics Delta Higher Institute of Engineering and Technology,Egypt
[6]Computer Science Department, Computer Science &Information Technology College, Anbar University, Anbar, Iraq
[7]Computer Engineering and Control Systems Department, Faculty of Engineering, Mansoura University Mansoura, Egypt
E-mail: mohammad.cs88@gmail.com[7]

Keywords: Data Poisoning, Adversarial Machine Learning, Training Data Security, Machine Learning Robustness, Machine Learning Vulnerabilities

Abstract

Data poisoning attacks, where adversaries manipulate training data to degrade model performance, are an emerging threat as machine learning becomes widely deployed in sensitive applications. This paper provides a comprehensive overview of data poisoning including attack techniques, adversary incentives, impacts on security and reliability, detection methods, defenses, and key research gaps. We examine label flipping, instance injection, backdoors, and other attack categories that enable malicious outcomes ranging from IP theft to accidents in autonomous systems. Promising detection approaches include statistical tests, robust learning, and forensics. However, significant challenges remain in translating academic defenses like adversarial training and sanitization into practical tools ready for operational use. With safety and trustworthiness at stake, more research on benchmarking evaluations, adaptive attacks, fundamental tradeoffs, and real-world deployment of defenses is urgently needed. Understanding vulnerabilities and developing resilient machine learning pipelines will only grow in importance as data integrity is fundamental to developing safe artificial intelligence.

I. Introduction

Machine learning (ML) has become a transformative technology across many industries [1], with models being embedded into sensitive systems like healthcare, finance, and transportation. However, the integrity and security of these models is directly dependent on the quality of the data used to train them. In recent years, the threat of data poisoning has emerged as a growing concern for ML security[2]. Data poisoning refers to the intentional corruption of training data to manipulate model behavior and degrade performance. By injecting crafted malicious data into the training set, attackers can introduce backdoors, increase error rates, induce misclassifications, or cause systems to crash once deployed. Unlike typical test data attacks, data poisoning enables adversaries to manipulate models prior to deployment by exploiting vulnerabilities in how models are trained[3]. While poisoning attacks were theoretically analyzed in early machine learning research, they have become an increasingly practical threat as large, complex models and massive training datasets grow more common. Real-world examples like Microsoft's Tay chatbot highlight the impacts of poisoning attacks that manipulate online training data. As machine learning expands into sensitive domains, data poisoning presents risks ranging from IP theft to catastrophic system failures. Addressing the data poisoning threat is challenging because it requires protecting the entire machine learning pipeline, from data collection through model development and deployment. Key research issues include developing poisoning detection mechanisms, designing robust training procedures, and formalizing realistic threat models. There remains a lack of benchmarks, best practices, and rigorous evaluation of data poisoning impacts across different learning algorithms, data modalities[3], and applications. This paper provides a comprehensive overview of the data poisoning threat, including methods of attack, motivations and incentives, impacts, analysis techniques, defenses, and critical gaps needing further research. We highlight key challenges that must be overcome to enable resilient and trustworthy machine learning in the face of adversarial data corruption. Understanding data poisoning vulnerabilities, their consequences, and how to mitigate them will only grow in importance as ML becomes even more pervasive[4].

II. Data Poisoning Attacks

Adversaries have developed an array of techniques to corrupt training data and achieve different malicious objectives. Common categories of data poisoning attacks include:

A. Label Flipping

This method involves changing the class labels for select instances in the training data. By flipping labels, attackers can manipulate the decision boundaries learned by classifiers. For example, spam email can be mislabeled as legitimate to avoid filters. Targeted label flipping aims to cause specific test instances to be misclassified[5].

https://doi.org/10.1049/icp.2024.0951

B. *Instance Injection*

The adversary generates new instances and injects them into the training data to skew the distribution. These injected instances consist of crafted feature vectors designed to introduce vulnerabilities. For example, attackers could insert noisy samples from the same class to degrade algorithm performance[6].

C. *Attribute Modification*

Instead of modifying labels, this attack tweaks input features/attributes of training instances. For example, slight pixel modifications to images can increase classification error rates. More advanced modifications can help evade defenses or mimic instances of a targeted class.

D. *Backdoor Insertion*

Backdoor attacks cause models to learn hidden patterns that serve as triggers to disrupt predictions at test time. For example, a specific pixel pattern in images could be linked to incorrect outputs. Backdoors provide covert control to attackers after deployment.

E. *Algorithm Subversion*

Rather than poisoning data itself, some techniques directly manipulate the training process. For example, adversaries could amplify weights for select instances, induce premature stopping, or incrementally poison across epochs.

These categories can impact machine learning pipelines in different ways depending on factors like attacker knowledge, capability, and objectives. Both untargeted and targeted poisoning attacks are possible, with the latter requiring more sophisticated manipulation. Understanding the diverse threat landscape is crucial for effective poisoning defense[7].

III. Threat Models and Adversary Incentives

To understand the risks posed by data poisoning, it is important to consider threat models—who are the attackers, what are their capabilities, and what are their motivations? Different adversaries are incentivized to corrupt training data based on their access, resources, and goals.

A. *Common adversaries include:*

Cybercriminals: Seeking financial gain, they may poison models used for spam detection, fraud analytics, recommendation systems.

Insiders: Employees or partners with data access may poison to sabotage systems or steal IP.

State-sponsored: Advanced threat actors aimed at disruption of services, surveillance, or IP theft.

Attackers can target components of the ML pipeline including data collection, labeling, model development, and deployment. Adversary knowledge levels range from no knowledge (black box) to full knowledge of datasets, model architectures, parameters, and defense mechanisms (white box)[1, 2, 5].

B. *Incentives for data poisoning include:*

Financial gain: Manipulate stock trades, product ratings and recommendations, ad placements

Service disruption: Degrade performance of online services, denial of service

IP theft: Extract proprietary algorithms, business logic, user data

Reputation damage: Create controversial outputs that damage brand image

Physical impacts: Cause accidents or dangerous conditions for autonomous systems

The high level of model dependence on training data combined with the breadth of incentives make ML an attractive attack surface. Defending against diverse data poisoning threats will require a combination of robust training procedures, anomaly detection, pipeline security, and monitoring during operations[8].

IV. Impact of Data Poisoning

Data poisoning attacks pose threats to the security, reliability, and integrity of machine learning systems. By manipulating training data, adversaries can cause a wide range of harms, depending on their incentives and capabilities. Some of the potential impacts include:

Security violations: Poisoning provides backdoors for adversaries to exploit, allows evasion of defenses, and enables IP or data theft.

Performance degradation: Higher error rates, reduced reliability, uncertainty in predictions due to skewed training distribution.

Algorithm subversion: Fundamental manipulation of model behavior contrary to original task objectives.

System crashes: Injection of maliciously crafted instances could cause failures or undefined behavior when deployed.

Data integrity loss: Training set no longer accurately represents real-world distribution after poisoning.

Financial impacts: Poisoning of trading, credit scoring, or ad auction models could enable fraud or stock manipulation.

Health and safety risks: Unreliable predictions for medical diagnosis, autonomous vehicle control systems could cause harm. The consequences reflect the broad adoption of machine learning in sensitive domains like healthcare, transportation, finance, defense, and more. Failures or security breaches in these systems due to data poisoning could lead to severe economic and physical harms.

Quantifying the potential scale of damages is challenging due to lack of experience with poisoning attacks in the real world. However, the fundamental dependence of ML on training data integrity necessitates development of robust defenses against poisoning given the incentives of attackers[9].

V. Detection of Data Poisoning Attacks

Detecting data poisoning attacks is critical for security, but poses significant technical challenges. Different detection methods have been proposed, but must balance tradeoffs between accuracy, robustness, and practicality. Common approaches include:

Statistical detection - Identify anomalies in data distribution, feature correlations, label patterns that indicate poisoning. Tests include density estimation, cluster analysis, dimensionality reduction techniques[10, 11].

Robust learning - Use algorithms like error amplification that increase influence of poisoning instances on model performance. The resultant degradation reveals presence of poisoning[12, 13].

Auditing - Rigorously verify provenance and lineage of training data. Audit logs, version control, and metadata inspection can reveal origins of corruption.

Forensics - Analyze models and data for evidence of backdoors, misclassification triggers, or other artifacts indicative of poisoning.

Challenges arise due to high dimensionality, concept drift in data streams, adaptive attackers that bypass detections, lack of labeled poisoning data, and computational bottlenecks. Statistical approaches struggle with large feature spaces. Robust learning often trades off accuracy for security. And auditing approaches face scalability issues across massive datasets[14, 15].

Ongoing research is focused on hybrid techniques, developing better statistical tests, and adopting methods from related domains like malware detection and network intrusion detection. More rigorous benchmarks on diverse poisoning attacks are needed to advance detection capabilities. Overall, a key priority is translating proposed techniques into practical mechanisms that account for real-world constraints.

VI. Defenses and Mitigation Approaches

A number of promising approaches have been proposed to defend against data poisoning threats, though many still remain theoretical or lack thorough evaluation. Defenses operate across the machine learning pipeline, aiming to prevent, detect, or reduce the influence of malicious data manipulation. Common categories include:

Data sanitization - Remove or repair parts of training data suspected of poisoning. Techniques include anomaly removal, noise reduction, and imputation of missing values.

Robust loss functions - Use modified losses like Reverse Cross-Entropy to reduce impact of poisoned points. Makes outlier manipulation more difficult.

Differential privacy - Restrict amount of influence any individual data instance can have on model. Provides stochasticity to mitigate poisoning impact.

Adversarial training - Improve regularization and intentionally inject small amounts of poison during training to increase model robustness.

Pipeline hardening - Apply access controls, provenance tracking, and integrity checks to training data collection and storage. Reduce attack surface area.

Monitoring - Deploy poisoning detectors during training and model performance monitors after deployment to identify anomalous behaviors.

However, most defenses involve some tradeoff between accuracy, robustness, and transparency. More research is needed to develop combined, adaptive solutions that remain practical for real-world systems. Efforts should focus on translating theoretic defenses into standardized toolkits, benchmarks, and best practices for practitioners.

VII. Gaps and Open Challenges

While an active area of research, defending against data poisoning attacks remains an open problem with many unanswered questions. Some of the key gaps and challenges going forward include:

Fundamental tradeoffs - There are inherent tensions between accuracy, robustness, and model transparency that must be navigated when dealing with poisoning. Formalizing these tradeoffs is an open theoretical problem[16-35].

Benchmarks and evaluation - Rigorous benchmarks with standardized datasets are needed to effectively evaluate performance of poisoning defenses across diverse threat models and learning algorithms.

Adaptive attacks - Attacks will continue to evolve and adapt to introduced defenses over time. Models for adversary incentives and costs are limited.

Transferability - The dynamics of how poisoning transfers across models, data distributions, and time is not well-understood. More research is needed on generalization.

Deployment challenges - Translating proposed academic defenses into real-world systems poses engineering, runtime, and monitoring hurdles that must be addressed.

Lack of tools and best practices - Practitioners need better tooling, techniques, and design guidelines to make models more resilient. Lack of experience deploying defenses remains a bottleneck[2, 6].

Addressing these gaps requires cross-disciplinary efforts combining security, machine learning, economics, and systems perspectives. Even with advances, the core dependence on training data necessitates a proactive, risk-aware approach to poisoning that evolves along with the threat landscape.

VIII. Discussion

This paper provided a comprehensive overview of the emerging threat of data poisoning attacks against machine learning systems. We discussed the diverse range of poisoning techniques adversaries employ to manipulate model behavior by corrupting training data. Beyond methods of attack, we analyzed different threat models and incentives that make machine learning an attractive target. Quantifying the impacts revealed serious risks to security, reliability, financial systems, and even human safety. Detecting poisoning attacks remains challenging, though promising directions include hybrid statistical and robust learning techniques. On the defense side, methods like sanitization, privacy mechanisms, adversarial training, and pipeline hardening show potential to mitigate data integrity threats. However, inherent tradeoffs imply a combination of approaches is likely needed.

Significant gaps remain when translating theoretical defenses into practical techniques that can be systematically implemented by practitioners. Lack of benchmarks, standardized datasets, and real-world experience deploying defenses hinder development of robust machine learning. Adaptive attacks that evolve to bypass defenses also pose an ongoing challenge.

Our analysis revealed that data poisoning represents an important threat going forward as reliance on machine learning grows across industries. However, with dedicated efforts to close gaps in detection, measurement, defenses, and adoption - the dangers of poisoning attacks can be reduced. Maintaining research momentum in this domain is crucial as data integrity is fundamental to developing trustworthy and safe AI systems.

IX. CONCLUSION

Data poisoning threatens the security foundations of modern machine learning systems. As models become more complex and training datasets grow larger, the opportunities for adversaries to manipulate behaviors through data corruption will only increase. However, the technical community has only scratched the surface in terms of rigorously analyzing vulnerabilities, quantifying real-world risks, and developing robust defenses.

This paper provided a comprehensive overview of the data poisoning landscape. We examined the diverse range of attack techniques, motivations, and impacts in depth. Promising detection and defense methods were highlighted, but significant gaps remain in translating academic research into practical tools ready for operational deployment. Adaptive attacks, benchmarking evaluations, fundamental tradeoffs, and lack of experience implementing defenses are key issues needing further research.

Data integrity is a prerequisite for safe and reliable machine learning. With so much at stake, the machine learning community must prioritize resilience against data poisoning threats. Maintaining research momentum and cross-disciplinary collaboration between security, ML researchers, and practitioners will be essential. The era of deep learning and big data presents risks, but also opportunities to develop more robust algorithms, training processes, and system architectures. With sustained effort, the threats introduced by data poisoning can be reduced to an acceptable level even as machine learning grows more pervasive across society.

References

[1] Z. Zhang et al., "Explainable Data Poison Attacks on Human Emotion Evaluation Systems Based on EEG Signals," (in English), IEEE Access, Article vol. 11, pp. 18134-18147, 2023.

[2] Z. Yang et al., "Data poisoning attacks against multimodal encoders," in International Conference on Machine Learning, 2023, pp. 39299-39313: PMLR.

[3] H. Bansal, N. Singhi, Y. Yang, F. Yin, A. Grover, and K.-W. Chang, "CleanCLIP: Mitigating Data Poisoning Attacks in Multimodal Contrastive Learning," arXiv preprint arXiv:2303.03323, 2023.

[4] M. Aljanabi and S. Y. Mohammed, "Metaverse: open possibilities," Iraqi Journal For Computer Science and Mathematics, vol. 4, no. 3, pp. 79-86, 2023.

[5] S. Wang, Q. Li, Z. Cui, J. Hou, and C. Huang, "Bandit-based data poisoning attack against federated learning for autonomous driving models," Expert Systems with Applications, vol. 227, p. 120295, 2023.

[6] Y. Li, J. Wang, and C. Wang, "Systematic Testing of the Data-Poisoning Robustness of KNN," in ACM SIGSOFT International Symposium on Software Testing and Analysis, 2023.

[7] J. Carnerero-Cano, L. Muñoz-González, P. Spencer, and E. C. Lupu, "Hyperparameter Learning under Data Poisoning: Analysis of the Influence of Regularization via Multiobjective Bilevel Optimization," arXiv preprint arXiv:2306.01613, 2023.

[8] H. Chen and F. Koushanfar, "Tutorial: Toward Robust Deep Learning against Poisoning Attacks," ACM Transactions on Embedded Computing Systems, vol. 22, no. 3, pp. 1-15, 2023.

[9] H. Lu, Z. Huang, and R. Wang, "Enhancing Healthcare Model Trustworthiness through Theoretically Guaranteed One-Hidden-Layer CNN Purification," in ICLR 2023 Workshop on Trustworthy Machine Learning for Healthcare, 2023.

[10] M. Al-Janabi and M. A. Ismail, "Improved intrusion detection algorithm based on TLBO and GA algorithms," Int. Arab J. Inf. Technol., vol. 18, no. 2, pp. 170-179, 2021.

[11] M. Aljanabi, M. A. Ismail, and V. Mezhuyev, "Improved TLBO-JAYA algorithm for subset feature selection and parameter optimisation in intrusion detection system," Complexity, vol. 2020, 2020.

[12] A. H. Ali, M. Aljanabi, and M. A. Ahmed, "Fuzzy generalized Hebbian algorithm for large-scale intrusion detection system," International Journal of Integrated Engineering, vol. 12, no. 1, pp. 81-90, 2020.

[13] M. Yaseen, H. S. Salih, M. Aljanabi, A. H. Ali, and S. A. Abed, "Improving Process Efficiency in Iraqi universities: a proposed management information system," Iraqi Journal For Computer Science and Mathematics, vol. 4, no. 1, pp. 211-219, 2023.

[14] S. Atheel Sabih, F. Y. Omar, A. Mohammad, Z. Abbood, and S. M. Mahdi, "SEEK Mobility Adaptive Protocol Destination Seeker Media Access Control Protocol for Mobile WSNs," Iraqi Journal For Computer Science and Mathematics, vol. 4, no. 1, pp. 130-145, 01/07 2023.

[15] S. Pal, R. Wang, Y. Yao, and S. Liu, "Towards Understanding How Self-training Tolerates Data Backdoor Poisoning," arXiv preprint arXiv:2301.08751, 2023.

[16] H. S. Salih, M. Ghazi, and M. Aljanabi, "Implementing an Automated Inventory Management System for Small and Medium-sized Enterprises," Iraqi Journal For Computer Science and Mathematics, vol. 4, no. 2, pp. 238-244, 2023.

[17] Makarovskikh, T., & Abotaleb, M. (2021). Comparison between two systems for forecasting Covid-19 infected cases. In Computer Science Protecting Human Society Against Epidemics: First IFIP TC 5 International Conference, ANTICOVID 2021, Virtual Event, June 28–29, 2021, Revised Selected Papers 1 (pp. 107-114). Springer International Publishing.

[18] Abotaleb, M. S. A., & Makarovskikh, T. (2021, July). The research of mathematical models for forecasting Covid-19 cases. In International Conference on Mathematical Optimization Theory and Operations Research (pp. 301-315). Cham: Springer International Publishing.

[19] Mishra, P., Yonar, A., Yonar, H., Kumari, B., Abotaleb, M., Das, S. S., & Patil, S. G. (2021). State of the art in total pulse production in major states of India using ARIMA techniques. Current Research in Food Science, 4, 800-806.

[20] Abotaleb, M. S., & Makarovskikh, T. (2021, September). Analysis of neural network and statistical models used for forecasting of a disease infection cases. In 2021 International Conference on Information Technology and Nanotechnology (ITNT) (pp. 1-7). IEEE.

[21] Sardar, I., Karakaya, K., Makarovskikh, T., Abotaleb, M., Aflake, S., & Mishra, P. (2021). MACHINE LEARNING-BASED COVID-19 FORECASTING: IMPACT ON PAKISTAN STOCK EXCHANGE. International Journal of Agricultural & Statistical Sciences, 17(1).

[22] Makarovskikh, T., & Abotaleb, M. (2022, May). Hyper-parameter tuning for long short-term memory (LSTM) algorithm to forecast a disease spreading. In 2022 VIII International Conference on Information Technology and Nanotechnology (ITNT) (pp. 1-6). IEEE.

[23] Makarovskikh, T., Salah, A., Badr, A., Kadi, A., Alkattan, H., & Abotaleb, M. (2022, May). Automatic classification Infectious disease X-ray images based on Deep learning Algorithms. In 2022 VIII International Conference on Information Technology and Nanotechnology (ITNT) (pp. 1-6). IEEE.

[24] Panyukov, A., Makarovskikh, T., & Abotaleb, M. (2022, September). Forecasting with Using Quasilinear Recurrence Equation. In International Conference on Optimization and Applications (pp. 183-195). Cham: Springer Nature Switzerland.

[25] Abotaleb, M. (2022). State of the art in wind speed in England using BATS. TBATS, Holt's Linear and ARIMA model, 1, 129-138.

[26] Alqahtani, F., Abotaleb, M., Kadi, A., Makarovskikh, T., Potoroko, I., Alakkari, K., & Badr, A. (2022). Hybrid deep learning algorithm for forecasting SARS-CoV-2 daily infections and death cases. Axioms, 11(11), 620.

[27] Alkanhel, R., El-kenawy, E. S., Abdelhamid, A., Ibrahim, A., Abotaleb, M., & Khafaga, D. (2023). Dipper Throated Optimization for Detecting Black-Hole Attacks in MANETs. Computers, Materials & Continua, 74, 1905–1921.

[28] Alkanhel, R., Khafaga, D. S., El-kenawy, E. S. M., Abdelhamid, A. A., Ibrahim, A., Amin, R., ... & El-den, B. M. (2023). Hybrid Grey Wolf and Dipper Throated Optimization in Network Intrusion Detection Systems. CMC-COMPUTERS MATERIALS & CONTINUA, 74(2), 2695-2709.

[29] Alkanhel, R., El-kenawy, E. S. M., Abdelhamid, A. A., Ibrahim, A., Alohali, M. A., Abotaleb, M., & Khafaga, D. S. (2023). Network Intrusion Detection Based on Feature Selection and Hybrid Metaheuristic Optimization. Computers, Materials & Continua, 74(2).

[30] Khafaga, D. S., El-kenawy, E. S. M., Karim, F. K., Abotaleb, M., Ibrahim, A., Abdelhamid, A. A., & Elsheweikh, D. L. (2023). Hybrid Dipper Throated and Grey Wolf Optimization for Feature Selection Applied to Life Benchmark Datasets. CMC-COMPUTERS MATERIALS & CONTINUA, 74(2), 4531-4545.

[31] El-kenawy, E. S. M., Abdelhamid, A. A., Ibrahim, A., Abotaleb, M., Makarovskikh, T., Alharbi, A. H., & Khafaga, D. S. Al-Biruni Earth Radius Optimization for COVID-19 Forecasting.

[32] El-kenawy, E. S. M., Abdelhamid, A. A., Alrowais, F., Abotaleb, M., Ibrahim, A., & Khafaga, D. S. (2023). Al-Biruni Based Optimization of Rainfall Forecasting in Ethiopia. Computer Systems Science & Engineering, 46(1).

[33] Makarovskikh, T., Panyukov, A., & Abotaleb, M. (2023, March). Monitoring and Forecasting Crop Yields. In International Conference on Parallel Computational Technologies (pp. 78-92). Cham: Springer Nature Switzerland.

[34] Makarovskikh, T., Panyukov, A., & Abotaleb, M. (2023, July). Using General Least Deviations Method for Forecasting of Crops Yields. In International Conference on Mathematical Optimization Theory and Operations Research (pp. 376-390). Cham: Springer Nature Switzerland.

[35] Alqahtani, F., Abotaleb, M., Subhi, A. A., El-Kenawy, E. S. M., Abdelhamid, A. A., Alakkari, K., ... & Kadi, A. (2023). A hybrid deep learning model for rainfall in the wetlands of southern Iraq. Modeling Earth Systems and Environment, 1-18.

[36]